

Module 1 Part 3
Concurrency Opportunity
Recognition -
K-means Algorithm

Carnegie Mellon University
18-645

Ian Lane
Jake Chong

18-645 – How to Write Fast Code? 1

What we discussed in Part 2:

- **Course Goal**
 - When your research/application needs to be fast, you will be able to:
 1. Feel comfortable hacking up a solution
 2. Leverage existing software building blocks
 3. Indicate which platform is the best one to use
 4. Reason about why a piece of existing code is slow
 5. Take care of potential performance bottlenecks
- **Hardware Architectures:**
 - Multicore vs Manycore – instruction latency vs throughput optimization
 - Opportunities in ILP, SIMD, SMP
 - Metrics for memory hierarchy
 - Metrics for system granularity

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2014 2

Questions you should ask yourself...

- What is the difference between concurrency and parallelism?
- What are the four key elements of the human problem solving process?
- What are the characteristics of a current algorithm implementation?
- What levels of concurrency can be **exposed** in the k-mean algorithm?
- What levels of parallelism are available to be **exploited**?
- What mapping between concurrency and parallelism can be **explored**?
- How is this relevant to writing fast code?

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

3

How to Write Fast Code?

Fast Platforms

- Multicore platforms
- Manycore platforms
- Cloud platforms



Good Techniques

- Data structures
- Algorithms
- Software Architecture

- This Lecture: Recognizing levels of concurrency in an application
- Effective Mapping of concurrency in an application with parallelism of a platform

→ **Fast code**

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

4

Outline

- The Application Developer
- Application-level Concurrency
- The Problem Solving Process

A number of figures in today's lecture were selected from Andrew Moore's tutorials in Statistical Data Mining (<http://www.cs.cmu.edu/~awm/tutorials>)

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

5

Distinction: Concurrency vs. Parallelism

Concurrency	Parallelism
The property of an application that... ...allows for tasks to have the potential to be... ...executed simultaneously	The property of a platform that... ...allows for tasks to have the potential to be... ...executed simultaneously
The application architecture in which... ... more than one task is active and able to... ... make progress at one time	The platform architecture in which.. ... more than one task can be active and... ... make progress at same time
We <u>expose</u> concurrency in our applications.	We <u>exploit</u> parallelism in our platforms.

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

6

The Application Developer

- Writing fast code is a process coherent with

“general problem solving behavior”

- Newell and Simon, Human Problem Solving (1972), pp. 72-73

- The process of problem solving involves:
 1. Understand the **current state**
 2. Observe the **internal representation**
 3. **Search** among alternatives
 4. Select from a set of **choices**

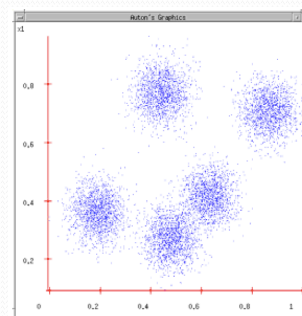
18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

7

k -means Problem

- Find k cluster centers that minimize the distance from each data point to a cluster center
- Important algorithm in machine learning:
 - Statistical data analysis
 - Vector quantization (Speech Recognition)
- NP-hard for arbitrary input
- **k -means algorithm** frequently finds a reasonable solutions quickly
- Issues:
 - Worst case running time is super-polynomial
 - Approximation can be arbitrarily bad



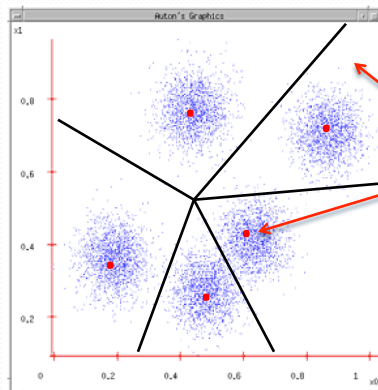
18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

8

k-means Problem

- Find k cluster centers that minimize the distance from each data point to a cluster center



Cluster

Cluster center (centroid)

k : Number of clusters (defined a-priori)
Cluster: Assignment of data points to a class
Cluster Center: μ of data points in a cluster

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

9

Related Problems and Algorithms

- k-means++**: Maximize scattering on initial cluster centers
- KD-trees**: Fast k -means - Pre-compute distance between data points
- x-means**: k -means with efficient estimation of the number of classes
- Gaussian Mixture Models:
 - Probabilistic assignments to clusters
 - Multivariate Gaussian distributions instead of means
- Expectation Maximization algorithms (EM algorithms)
 - Find maximum likelihood estimates of parameters in a statistical model, where the model depends on unobserved latent variables.
- Expectation Maximization Algorithms for Conditional Likelihoods
 - Estimate parameters in a statistical model to optimize conditional likelihood (where the objective function is a rational function)

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

10

k-means Algorithm (“Lloyd’s algorithm”)

- Given an initial set of k means $\mathbf{m}_1^{(1)}, \dots, \mathbf{m}_k^{(1)}$
- Expectation Step:** Assign each observation to the cluster with the closest mean

$$S_i^{(t)} = \left\{ \mathbf{x}_j : \|\mathbf{x}_j - \mathbf{m}_i^{(t)}\| \leq \|\mathbf{x}_j - \mathbf{m}_{i^*}^{(t)}\| \text{ for all } i^* = 1, \dots, k \right\}$$

- Maximization Step:** Calculate the new means to be the centroid of the observations in the cluster.

$$\mathbf{m}_i^{(t+1)} = \frac{1}{|S_i^{(t)}|} \sum_{\mathbf{x}_j \in S_i^{(t)}} \mathbf{x}_j$$

- Iterate until convergence or stopping criteria met

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

11

The Algorithm

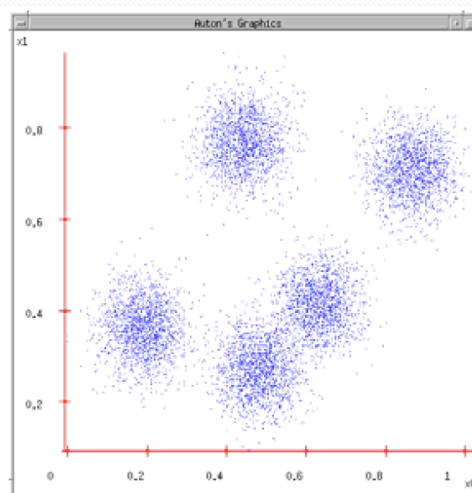
Example:

$k=5$

Distance metric=euclidean

Dimensions=2

1. Randomly select k cluster Centers
2. Assign closest Center to each data point
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met



18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

12

The Algorithm

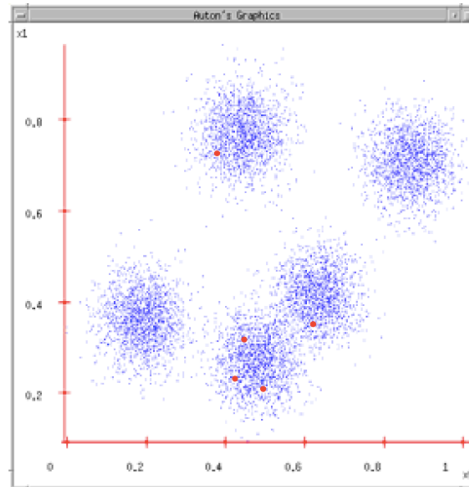
Example:

$k=5$

Distance metric=euclidean

Dimensions=2

1. Randomly select k cluster Centers
2. Assign closest Center to each data point
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met



18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

13

The Algorithm

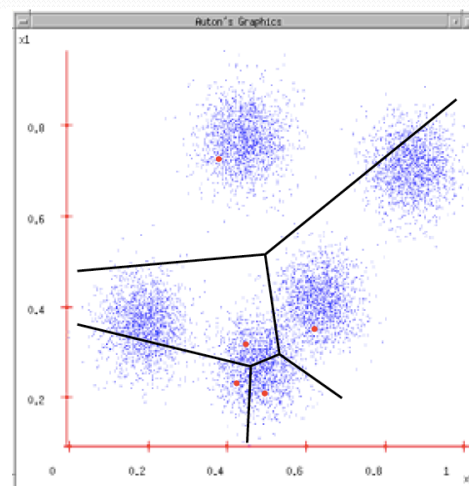
Example:

$k=5$

Distance metric=euclidean

Dimensions=2

1. Randomly select k cluster Centers
2. Assign each data point to closest Center
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met



18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

14

The Algorithm

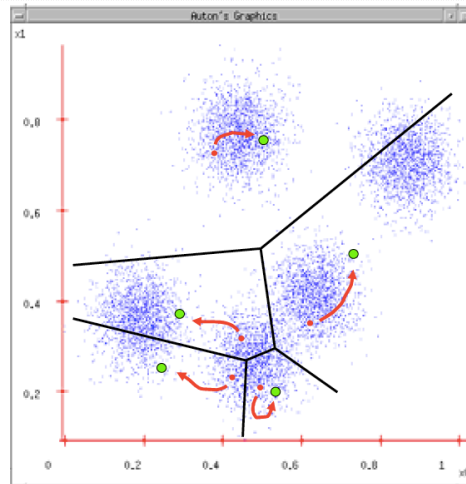
Example:

$k=5$

Distance metric=euclidean

Dimensions=2

1. Randomly select k cluster Centers
2. Assign closest Center to each data point
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met



18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

15

The Algorithm

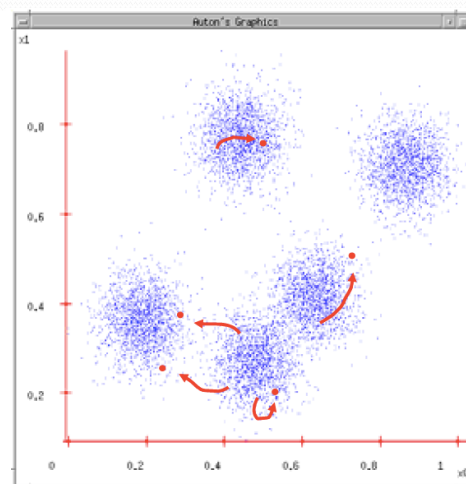
Example:

$k=5$

Distance metric=euclidean

Dimensions=2

1. Randomly select k cluster Centers
2. Assign closest Center to each data point
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met



18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

16

The Phases

1. **Initialization:** Randomly select k cluster centers
 - Select k samples from data as initial centers [Forgy Partition]
2. **Expectation:** Assign each data point go closest center
 - Compare each data point (N) to each cluster center (k)
 - Distance Metric: Euclidean distance (D dimensions)
3. **Maximization:** Update centers based on assignments
 - For each cluster (k) compute mean (D dimensions) from data points assigned to that cluster
4. **Evaluate:** Re-iterate steps 2-3 until convergence or stopping criteria met
 - Percentage of data points re-assigned
 - Number of iterations (2-3)

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

17

A Fast Implementation of k -means

- Following the process of problem solving with k -means:

1. Understand the **current state**

- Running on a platform
- Using a specific set of resources
- Achieving a specific performance
- Meeting a specific criteria/requirement

2. Observe the **internal representation**

3. **Search** among alternatives

4. Select from a set of **choices**

Assumption:

Starting from a functionally correct reference implementation

Implication:

Must observe the *current state* and *implementation requirements* before starting to solve a problem

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

18

Understanding the Current State

- **Running on a platform**
 - Platform: *Linux + GCC on x86 multicore processor*
- **Using a specific set of resources (ghcXX)**
 - Computation: *2 to 8 cores, 2 to 8 way SIMD*
 - Data: *32KB L1, 256KB L2, shared L3 cache, 2 to 16GB DRAM*
 - Synchronization: *on-chip shared-memory abstraction*
- **Achieving a specific performance**
 - *As measured in Homework 1*
- **Meeting a specific criteria/requirement**
 - Matrix-Multiply: *5x performance on largest size test set*
 - **k-means:** **1.5x performance on largest set**

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

19

What to Measure for Performance?

- Lecture coming up:
Module 2 Part 3- **Performance Analysis: Roofline Model**

- Before that, a few simple techniques:
 - Observe the phases of execution
 - Characterize the execution time break downs
 - Reason about why a piece of code is slow
 - Identify performance bottlenecks

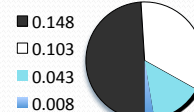
RTF: Real Time Factor (Proc Time / Real Time)

■ Phase 1 ■ Phase 3
□ Phase 2 ■ Seq. Overhead

Sequential Implementation



Parallel Implementation



Kisun You, Jike Chong, Youngmin Yi, Ekaterina Gonina, Christopher Hughes, Yen-Kuang Chen, Wonyong Sung, Kurt Keutzer, "Parallel Scalability in Speech Recognition: Inference engine in large vocabulary continuous speech recognition", IEEE Signal Processing Magazine, vol. 26, no. 6, pp. 124-135, November 2009.

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

20

Current state: k -means algorithm

- 4 Phases (Initialization, Expectation, Maximization, Evaluate)
 - Majority of time spent on Expectation and Maximization phases
- Entire data set can fit in memory on a single machine
- Number of samples (N) and feature dimensions (D) vary significantly
- Evaluation for any number of clusters ($2 \leq k \leq 300$)
- Example Data Sets:
 - *ionosphere_scale*: 351 Samples, 34 Dimensions
 - *svmguide*: 7089 Samples, 4 Dimensions
 - *cod-rna*: 59535 Samples, 8 Dimensions
 - *ljcnn1*: 191681 Samples, 22 Dimensions

← Grade for mini-project1 base

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

21

A Fast Implementation of k -means

- Following the process of problem solving with k -means:
 1. Understand the **current state**
 2. Observe the **internal representation**
 - Application structure
 - Identified four phases of execution
 - Implementation concerns
 - Task considerations
 - Data representations
 - Concurrency opportunities
 3. **Search** among alternatives
 4. Select from a set of **choices**

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

22

Example Code (Initialize)

kmeans/seq_kmeans.c
lines: 116-119

```
....  
/* pick first numClusters elements of objects[] as initial cluster centers*/  
for (i=0; i<numClusters; i++)  
    for (j=0; j<numCoords; j++)  
        clusters[i][j] = objects[i][j];  
....  
  
__inline static float euclid_dist_2(int numdims, float *coord1, float *coord2)  
{  
    int i;  
    float ans=0.0;  
    for (i=0; i<numdims; i++)  
        ans += (coord1[i]-coord2[i]) * (coord1[i]-coord2[i]);  
    return(ans);  
}
```

Define distance metric (Euclidian)
kmeans/seq_kmeans.c
lines: 51-63

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2014 23

Example Code (Initialize)

kmeans/seq_kmeans.c
lines: 116-119

Active Data Structures

objects: $N \times D$ clusters: $k \times D$

```
....  
/* pick first numClusters elements of objects[] as initial cluster centers*/  
....  
  
__inline static float euclid_dist_2(int numdims, float *coord1, float *coord2)  
{  
    Active Data Structures  
    coord1:  $1 \times D$   
    coord2:  $1 \times D$   
    ans: 1  
    D (sum reduction) ← euclid_dist_2()  
}
```

Define distance metric (Euclidian)
kmeans/seq_kmeans.c
lines: 51-63

Possible concurrencies:
D (sum reduction) ← euclid_dist_2()

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2014 24

Example Code (Expectation)

kmeans/seq_kmeans.c
lines: 136-147

```
....
delta = 0.0;
for (i=0; i<numObjs; i++) {
    /* find the array index of nearest cluster center */
    index = find_nearest_cluster(numClusters, numCoords, objects[i], clusters);

    ...

    /* if membership changes, increase delta by 1 */
    if (membership[i] != index) delta += 1.0;

    /* assign the membership to object i */
    membership[i] = index;
}
...
```

Evaluate distance to each cluster centroid and select closest

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2014 25

Example Code (Expectation)

kmeans/seq_kmeans.c
lines: 136-147

```
....
delta = 0.0;
for (i=0; i<numObjs; i++) {
    /* find the array index of nearest cluster center */
    index = find_nearest_cluster(numClusters, numCoords, objects[i], clusters);

    ...

    membership[i] = index;
}
...
```

Active Data Structures

objects: $N \times D$ clusters: $k \times D$ membership: $N \times 1$

Possible Concurrencies:

- N (independent)
- D (sum reduction) \leftarrow euclid_dist_2()
- k (min reduction)

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2014 26

Example Code (Maximization)

kmeans/seq_kmeans.c
lines: 148-162

```
....  
/* update new cluster centers : sum of objects located within */  
newClusterSize[index]++;  
for (j=0; j<numCoords; j++)  
    newClusters[index][j] += objects[i][j];  
}  
  
/* average the sum and replace old cluster centers with newClusters */  
for (i=0; i<numClusters; i++) {  
    for (j=0; j<numCoords; j++) {  
        if (newClusterSize[i] > 0)  
            clusters[i][j] = newClusters[i][j] / newClusterSize[i];  
        newClusters[i][j] = 0.0; /* set back to 0 */  
    }  
    newClusterSize[i] = 0; /* set back to 0 */  
}
```

prepare for next iteration

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2014 27

Example Code (Maximization)

kmeans/seq_kmeans.c
lines: 148-162

```
....  
/* update new cluster centers : sum of objects located within */  
newClusterSize[index]++;  
for (j=0; j<numCoords; j++)  
    newClusters[index][j] += objects[i][j];  
}
```

Active Data Structures

objects:	$N \times D$	membership:	$N \times 1$	NewClusters:	$k \times D$
----------	--------------	-------------	--------------	--------------	--------------

Possible Concurrencies:
D (independent)
N (Histogram computation into k bins)

newClusterSize[i] = 0; /* set back to 0 */
}

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2014 28

Example Code (Evaluate)

kmeans/seq_kmeans.c
lines: 164-165

```
....  
    delta /= numObjs;  
    } while (delta > threshold && loop++ < 500);  
...
```

Two stopping criteria
defined

- Two stopping criteria:
 - Percentage of data points that change class < *threshold*
 - Max 500 iterations (Assign and Update)

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

29

The phases - concurrency

1. **Initialization:** Randomly select k cluster centers
2. **Expectation:** Assign closest center to each data point
 - N (independent)
 - k (min reduction)
 - D (sum reduction)
3. **Maximization:** Update centers based on assignments
 - D (independent)
 - N (Histogram computation into k bins)
4. **Evaluate:** Re-iterate steps 2-3 until convergence

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

30

A Fast Implementation of k -means

- Following the process of problem solving with k -means:
 1. Understand the **current state**
 2. Observe the **internal representation**
 3. **Search** among alternatives
 4. Select from a set of **choices**

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

31

Search Among Alternatives

- Given the observed internal representations and the concurrency opportunities...
- What are the implementation **alternatives**?
 - Different mapping of **application concurrency** to **platform parallelism**
- The **search** process
 - More complex than one **application concurrency** to one **platform parallelism**
 - May want to sequentialize some operations:
 - Some parallel operations are as “**work-efficient**” as sequential operations
 - Reduction – sequential: $O(N)$, Parallel: $O(N \log N)$
 - One level of concurrency could map to multiple levels of parallelism

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

32

Multicore and Manycore Parallelism

- Similar in scaling trends:
 - Increasing vector unit width
 - Increasing numbers of cores per die
 - Increasing bandwidth to off-chip memory
- Different in optimization points

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2014 33

Memory Hierarchy

- Cache:**
 - A piece of fast memory close to the compute modules in a microprocessor
 - Allows faster access to a limited amount of data
 - Physical properties of wires and transistors determines trade-off between **cache capacity** and **access throughput/latency**

Memory Level	Capacity	Throughput	Latency
Network	Exa-Bytes	0.1 GB/s	~1-10B cycles
Disk	3TB	6 GB/s	~1M cycles
Memory	16GB	34.08 GB/s	~200 cycles
L3 cache	8 MB	108.8 GB/s	26-31 cycles
L2 cache	256 KB	108.8 GB/s	12 cycles
L1 cache	32 KB Data Cache	163.2 GB/s	4-6 cycles
core	Intel Core2 – 2600k @ 3.4GHz	(viewed from one core)	

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2014 34

Mapping Concurrency to Parallelism

- How does it map to the platform?
 - SIMD level parallelism
 - Core level parallelism
- How does it map to the cache hierarchy?
 - What data is required for each concurrent operation?
 - What are the synchronization points in the algorithm?
- Expectation & Maximization Phases
 - SIMD & core-level parallelism across data-points (N)
 - Update membership for each data point sequentially
 - Compute distance to each cluster center and select index with min. distance
 - Histogram computation (summation / assignment count for new clusters)
 - Other possible concurrency mappings?

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

35

A Fast Implementation of k -means

- Following the process of problem solving with k -means:
 1. Understand the **current state**
 2. Observe the **internal representation**
 3. **Search** among alternatives
 4. Select from a set of **choices**
 - Does solution met required criteria
 - How to evaluate a mapping?
 - Efficiency: Runs quickly, makes good use of computational resources
 - Simplicity: Easy to understand code is easier to develop, debug, verify and modify
 - Portability: Should run on widest range of parallel computers
 - Scalability: Should be effective on a wide range of processing elements
 - Other considerations: Practicality, Hardware, Engineering cost

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

36

Evaluate Choice

- Expectation & Maximization Phases
 - SIMD & core-level parallelism across data-points (N)
 - Update membership for each data point sequentially
 - Histogram computation (summation / assignment count for new clusters)
- OpenMP
- How we can evaluate the choice and make a decision
 - Efficiency
 - Simplicity / Maintainability
 - Portability
 - Scalability

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

37

How to write fast code

- **Expose** concurrencies in applications and algorithms
 - **Module 1 Part 3:** “Concurrency Opportunity Recognition”
 - Mini-Projects (1-3) & Term Project
- **Exploit** parallelisms on application platform
 - **Module 1 Part 2:** “Advanced Parallel Hardware Architectures”
 - Mini-Projects (1-3) & Term Project
- **Explore** mapping between concurrency and parallelism
 - The rest of the semester....
 - Abstractions to support mapping of concurrencies to parallelisms
 - OpenMP [Module 2]
 - CUDA [Module 3]
 - Hadoop [Module 4]

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2014

38