

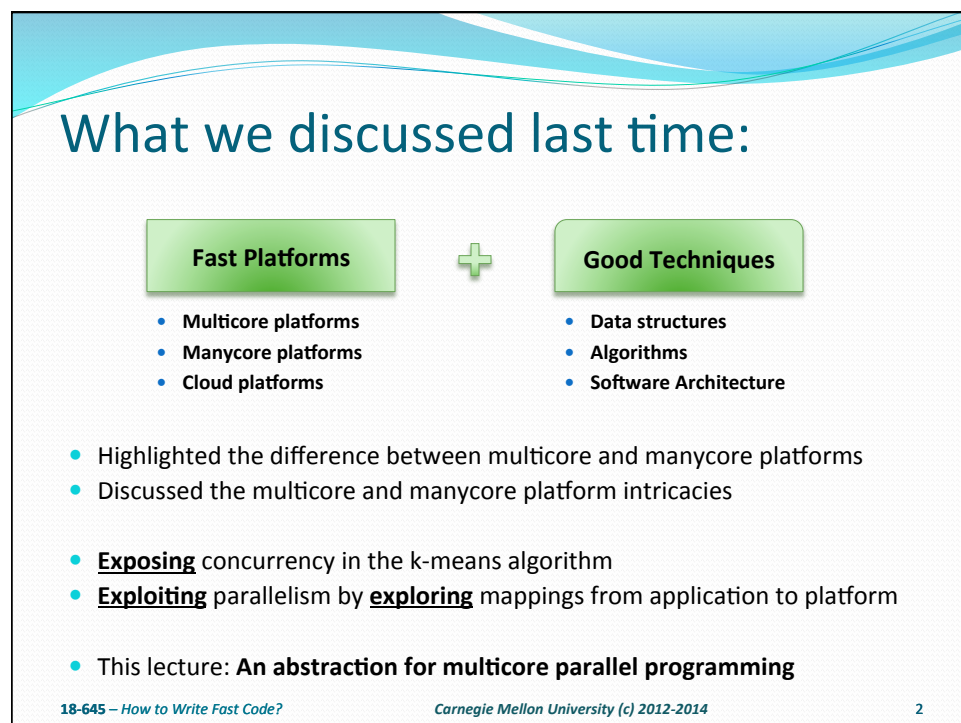
Module 2 Part 1

Application Design for Multicore Programming

Carnegie Mellon University
18-645

Jike Chong
Ian Lane

18-645 – How to Write Fast Code? 1



What we discussed last time:

Fast Platforms

- Multicore platforms
- Manycore platforms
- Cloud platforms

+

Good Techniques

- Data structures
- Algorithms
- Software Architecture

- Highlighted the difference between multicore and manycore platforms
- Discussed the multicore and manycore platform intricacies
- **Exposing** concurrency in the k-means algorithm
- **Exploiting** parallelism by **exploring** mappings from application to platform
- This lecture: **An abstraction for multicore parallel programming**

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2012-2014 2

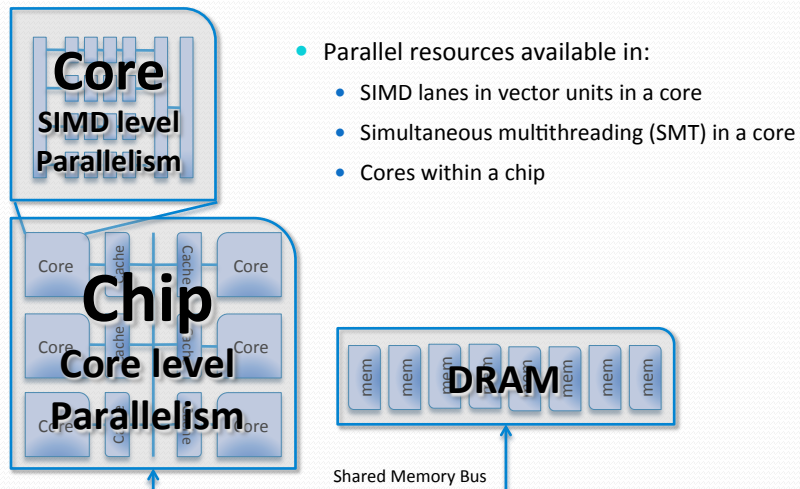
Outline

- Multicore Shared-Memory Platforms
- OpenMP – API for Shared-Memory Multi-processing
 - Overview
 - Example: numerical integration of Pi
 - Shared variable synchronization
 - SPMD vs worksharing
 - Data environment options
 - Advanced worksharing options
 - Synchronization options
- How is This Relevant to Writing Fast Code?

Answers you should know after this...

- What are the exploitable levels of parallelism in a multicore processor?
- What is SPMD? And how to use OpenMP to do SPMD?
- What's the difference between "critical" and "atomic"?
- How to reduce synchronization cost and avoid "false sharing"?
- What are the scheduling, reduction, data sharing, and synchronization options for OpenMP?
- How is this relevant to writing fast code?

Multicore Shared-Memory Platforms




18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

5


Exploiting Different Levels of Parallelism

- How do we write fast code to exploit the many levels of parallelism on a multicore platform?



SIMD-Level Parallelism

Exploited using vectorizing compiler and hand-code intrinsics



SMT-Level Parallelism

OS abstract it to core-level parallelism

Core-Level Parallelism

Using **threads** to describe work done on different cores

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

6

What is a thread?

- A thread of execution is a unit of processing scheduled by the OS
 - Allows system resources to be shared efficiently

Processes	Threads	Fiber
Exist independently in the OS	Subsets of processes	Subsets of threads
Independent states: virtual memory space, handles to system objects, a security context, a unique process identifier, environment variables, a priority class, minimum and maximum working set sizes, and at least one thread of execution	Threads in a process share: virtual address space and system resources Independent: exception handlers, a scheduling priority, a unique thread identifier, and thread context	Fibers within a thread share: virtual address space, system resources, thread context, a manually specified schedule Independent: exception handlers, a unique fiber identifier
Interact through system-provided inter-process communication mechanisms	Interact through shared memory	Interact through shared memory
OS Preemptive	OS Preemptive	Non-Preemptive

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

7

Landscape of Thread Programming

- POSIX threads
 - POSIX: Portable Operating System Interface for Unix – IEEE Standard
 - Defines a set of C programming language types, functions and constants
 - **Manually expose and manage all concurrency with the API**
- OpenMP
 - OpenMP: Open Multi-Processing – an API for shared memory multiprocessing
 - **Programmer hints at the concurrency**
 - **Compiler manages the parallelism**
- Intel Cilk Plus, Intel Thread Building Block (TBB):
 - Fork-join parallelism (“spawn” and “sync” in Cilk, Parallel loops/containers in TBB)
 - **Programmer exposes the concurrency**
 - **Runtime determines what is run in parallel**

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

8

Outline

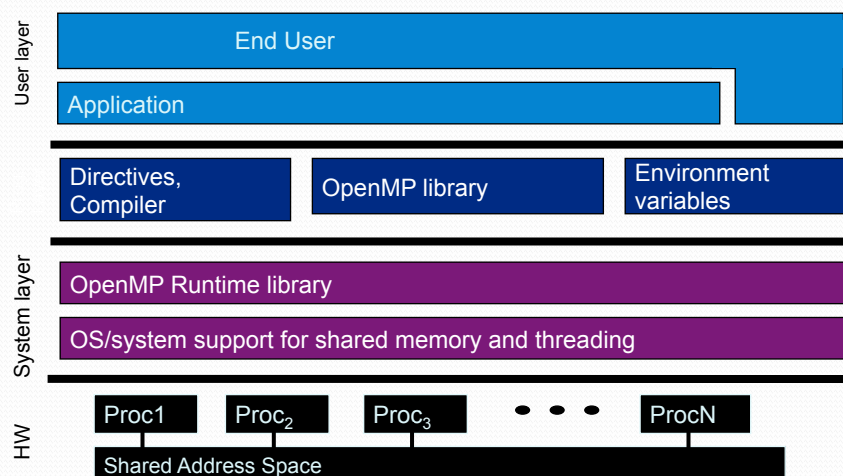
- Mutlicore Shared-Memory Platforms
- OpenMP – API for Shared-Memory Multi-processing
 - **Overview**
 - Example: numerical integration of Pi
 - Shared variable synchronization
 - SPMD vs worksharing
 - Data environment options
 - Advanced worksharing options
 - Synchronization options
- How is This Relevant to Writing Fast Code?

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

9

OpenMP Basic Defs: Solution Stack



Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

10

Scope of discussion for OpenMP

Topic	Concepts
I. OMP Intro	Parallel regions
II. Creating threads	Parallel, default data environment, runtime library calls
III. Synchronization	False sharing, critical, atomic
IV. Parallel loops	For, schedule, reduction
V. Odds and ends	Single, sections, master, runtime libraries, environment variables, synchronization, etc.
VI. Data Environment	Data environment details, software optimization
VII. OpenMP 3 and tasks	Tasks and other OpenMP 3 features
VIII. Memory model	The need for flush, but its actual use is left to an appendix
IX. Threadprivate	Modular software

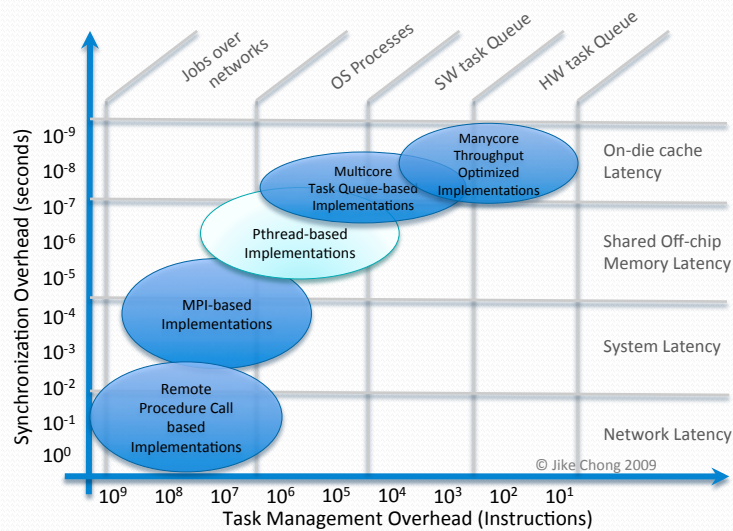
Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

11

Spectrum of Granularity



18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

12

Example 1: Hello World

Verifying that your OpenMP environment works

- Write a multithreaded program where **each thread** prints “hello world”.

```
void main()
{

    int ID = 0;
    printf(" hello(%d) ", ID);
    printf(" world(%d) \n", ID);

}
```

Mattson et al, A “Hands-on” Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

13

Example 1: SPMD with “omp parallel”

Verifying that your OpenMP environment works

- Write a multithreaded program where each thread prints “hello world”.

```
#include <omp.h>
void main()
{

    #pragma omp parallel
    {

        int ID = omp_get_thread_num();
        printf(" hello(%d) ", ID);
        printf(" world(%d) \n", ID);

    }

}
```

Switches for compiling and linking

```
gcc -fopenmp gcc
pgcc -mp pgi
icl /Qopenmp intel (windows)
icc -openmp intel (linux)
```

Mattson et al, A “Hands-on” Introduction to OpenMP

18-645 – How to Write Fast Code?

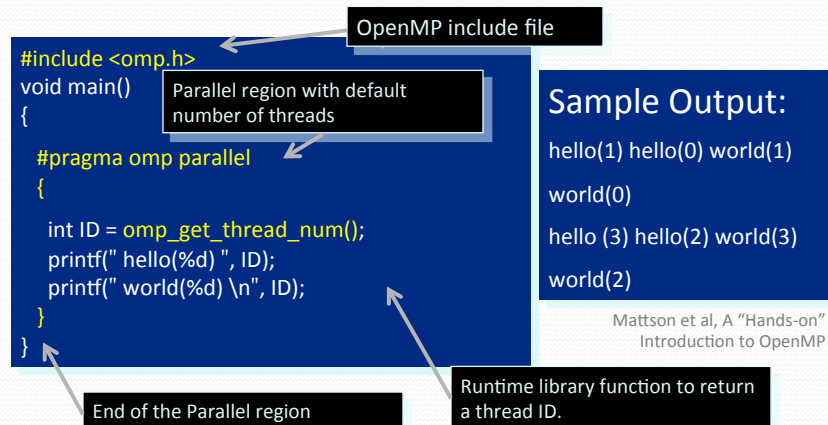
Carnegie Mellon University (c) 2012-2014

14

Example 1: SPMD with “omp parallel”

A multi-threaded “Hello world” program

- Write a multithreaded program where each thread prints “hello world”.



18-645 – How to Write Fast Code?

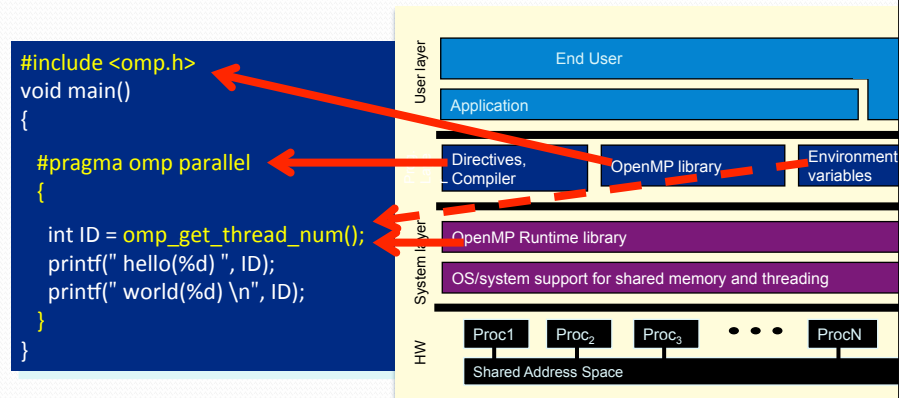
Carnegie Mellon University (c) 2012-2014

15

Example 1: SPMD with “omp parallel”

A multi-threaded “Hello world” program

- Write a multithreaded program where each thread prints “hello world”.




18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

16

OpenMP: Execution Pattern



- **Managed Fork-Join Parallelism:**
 - **Master thread** creates a **team of threads** as needed
 - Parallelism added incrementally until performance goals are met
 - i.e. the **sequential program evolves into a parallel program**

Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

17

Thread Creation: Parallel Regions

- Threads in OpenMP are created with the parallel construct
 - To create a 4 thread Parallel region:

Each thread executes a copy of the code within the structured block

```
double A[1000];
omp_set_num_threads(4);
#pragma omp parallel
{
    int ID = omp_get_thread_num();
    foo(ID,A);
}
```

Runtime function to request a certain number of threads

Runtime function returning a thread ID

Each thread calls foo(ID,A) for ID = 0 to 3

Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

18

Thread Creation: Parallel Regions

- Threads in OpenMP are created with the parallel construct
 - To create a 4 thread Parallel region:

Each thread executes a copy of the code within the structured block

```
double A[1000];
#pragma omp parallel num_threads(4)
{
    int ID = omp_get_thread_num();
    foo(ID,A);
}
```

Runtime function to request a certain number of threads

Runtime function returning a thread ID

Each thread calls foo(ID,A) for ID = 0 to 3

Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

19

Outline

- Mutlicore Shared-Memory Platforms
- OpenMP – API for Shared-Memory Multi-processing
 - Overview
 - **Example: numerical integration of Pi**
 - Shared variable synchronization
 - SPMD vs worksharing
 - Data environment options
 - Advanced worksharing options
 - Synchronization options
- How is This Relevant to Writing Fast Code?

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

20

Example: Numerical Integration of Pi


Mathematically, we know that:

$$\int_0^1 \frac{4}{1+x^2} dx = 4 \cdot \tan^{-1} x \Big|_0^1 = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i .



Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

21

Numerical Integration of Pi

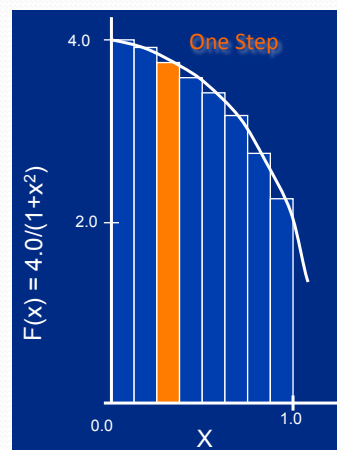
```
#include <stdio.h>
static long num_steps = 100000; double step;

int main ()
{
    int i; double pi, sum;
    step = 1.0/(double) num_steps;

    int i; double x;

    for (i=0, sum=0.0; i< num_steps; i++){
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
    }

    pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```



Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

22

Numerical Integration of Pi

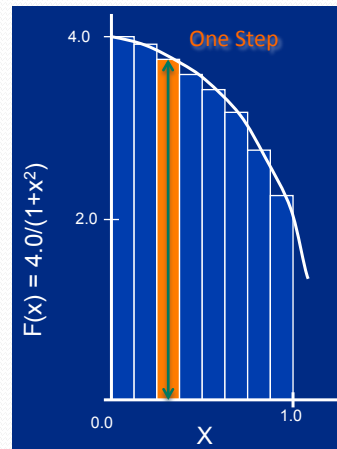
```
#include <stdio.h>
static long num_steps = 100000; double step;

int main ()
{
    int i; double pi, sum;
    step = 1.0/(double) num_steps;

    int i; double x;

    for (i=0, sum=0.0; i< num_steps; i++){
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
    }

    pi = sum * step;
    printf("pi = %i\n", pi);
    return 0;
}
```



Mattson et al, A "Hands-on" Introduction to OpenMP


18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

23

Concurrency Opportunities

- What is parallelizable?
 - Calculation of each step is independent
- One can map each of the steps to an independent thread
 - **What's wrong with this approach?**
- Significant thread management overhead
 - Compute each step as a separate thread involves significant thread management overhead (hundreds of cycles)



Mattson et al, A "Hands-on" Introduction to OpenMP



18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

24

Concurrency Opportunities

- What is parallelizable?
 - Calculation of each step is independent
- Alternative:
 - Specify a fixed set of parallel entities e.g. two threads (red and yellow threads)
 - Parallelize workload over the available number of threads





18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

25

Implementation View

- Work is evenly distributed across two processors



18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

26

Example 2: Numerical Integration of Pi

```
#include <stdio.h>
static long num_steps = 100000000; double step;


int main ()
{
    int i; double pi, sum = 0.0;
    step = 1.0/(double) num_steps;

    double x;

    for (i=0; i< num_steps; i++){
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
    }

    pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds



Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

27

Outline

- Mutlicore Shared-Memory Platforms
- OpenMP – API for Shared-Memory Multi-processing
 - Overview
 - Example: numerical integration of Pi
 - **Shared variable synchronization**
 - SPMD vs worksharing
 - Data environment options
 - Advanced worksharing options
 - Synchronization options
- How is This Relevant to Writing Fast Code?

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

28

Example 3: Critical Region

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double pi; int nthreads; double sum=0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel num_threads(NUM_THREADS)
    {
        int i, id = omp_get_thread_num(); double x;
        int nthrds = omp_get_num_threads();
        for (i=id; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            #pragma omp critical
            {sum += 4.0/(1.0+x*x);}
        }
        if (id == 0)    nthreads = nthrds;
    }
    pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds



64.45 seconds

Why so slow?

Anything in the critical region, OpenMP would make sure it is sequentialized.

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

29

Example 4: Atomic Operation

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double pi; int nthreads; double sum=0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel num_threads(NUM_THREADS)
    {
        int i, id = omp_get_thread_num(); double x, tmp;
        int nthrds = omp_get_num_threads();
        for (i=id; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            tmp = 4.0/(1.0+x*x);
            #pragma omp atomic
            {sum += tmp;}
        }
        if (id == 0)    nthreads = nthrds;
    }
    pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds



4.93 seconds

Better, why still so slow?


18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

30

Shared Memory Model

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP



- All threads share an address space
- Multiple copies of data may be present in various levels of cache

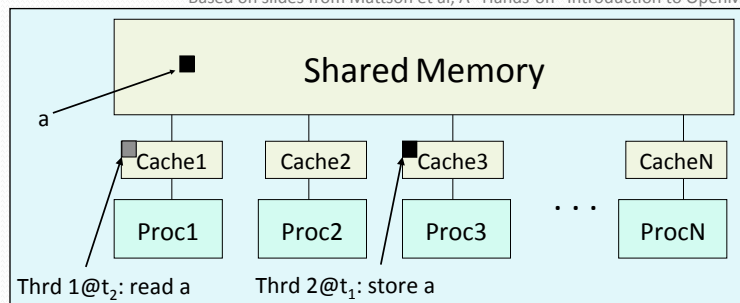
18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

31

Shared Memory Model

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP



- A piece of data (at address "a") can be:
 - Written into memory by Thrd 2 at t_1
 - Then, read by Thrd 1 at t_2
- We say the data at address "a" is in a **shared** state

Mattson et al, A "Hands-on" Introduction to OpenMP

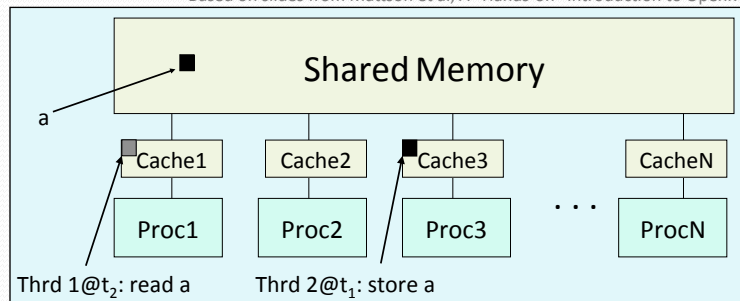
18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

32

Managing Coherence Costly

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

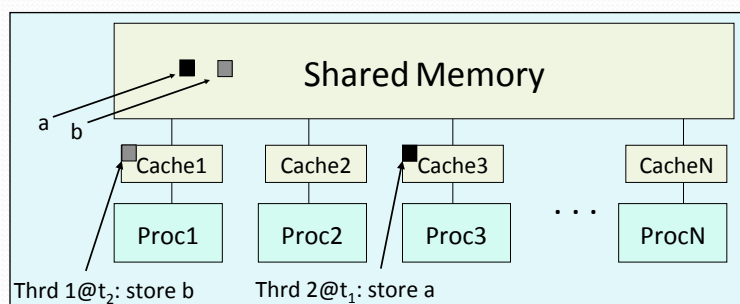


- Synchronization protocols are implemented in hardware to manage coherence between caches
 - e.g. **MESI Protocol** (Papamarcos & Patel 1984)
 - Allows data to be in one of four states: *Modified, Exclusive, Shared, Invalid*
- Transparent to software, but could have **severe performance penalties**

18-645 – How to Write Fast Code?

Mattson et al, A "Hands-on" Introduction to OpenMP
Carnegie Mellon University (c) 2012-2014 33

Threads can Write to Different Addresses...



- Avoid cache **synchronization protocols** induced **severe performance penalties** when no sharing is required
- Technique:
 - Write to separate addresses during parallel computation

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

34

Example 5: Separate Result Storage

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double pi; int nthreads; double sum[NUM_THREADS]
    step = 1.0/(double) num_steps;
    #pragma omp parallel num_threads(NUM_THREADS)
    {
        int i, id = omp_get_thread_num(); double x;
        int nthrds = omp_get_num_threads();
        for (i=id, sum[id]=0.0; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            sum[id] += 4.0/(1.0+x*x);
        }
        if (id == 0)    nthreads = nthrds;
    }
    for(i=0, pi=0.0; i<nthreads; i++) pi += sum[i] * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds

1.47 seconds

Better, but
why slower
2x speedup?

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

35

Structure of a Cache

- **Principle of Locality**
 - The phenomenon of the same value or related storage locations being frequently accessed, usually amenable to performance optimization
- **Temporal Locality**
 - Reuse of specific data and/or resources within relatively small time durations
- **Spatial Locality**
 - Use of data elements within relatively close storage locations

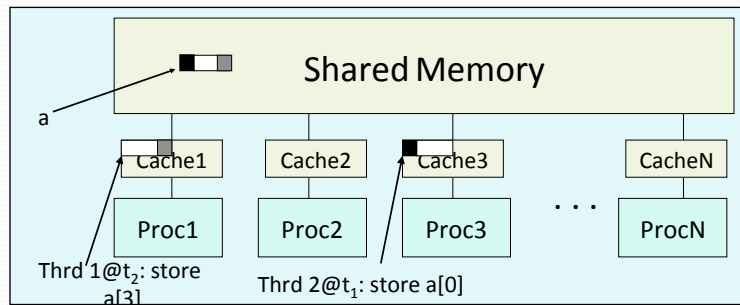
$$\begin{array}{c}
 \begin{array}{ccc}
 A & & B \\
 \begin{pmatrix} 4 & 5 & -5 \\ -1 & -4 & -2 \\ -3 & 1 & 5 \\ 2 & 1 & 4 \end{pmatrix} & & \begin{pmatrix} -4 & 4 & -1 & 3 & -4 \\ -1 & -5 & -5 & 4 & -5 \\ -1 & -1 & 0 & -1 & 1 \end{pmatrix} \\
 \\
 AB = \begin{pmatrix} -16 & -4 & -29 & 37 & -46 \\ 10 & 18 & 21 & -17 & 22 \\ 6 & -22 & -2 & -10 & 12 \\ -13 & -1 & -7 & 6 & -9 \end{pmatrix} \\
 \\
 \begin{pmatrix} -1 \end{pmatrix} \begin{pmatrix} 3 \end{pmatrix} + \begin{pmatrix} -4 \end{pmatrix} \begin{pmatrix} 4 \end{pmatrix} + \begin{pmatrix} -2 \end{pmatrix} \begin{pmatrix} -1 \end{pmatrix} \\
 -3 \quad + \quad -16 \quad + \quad 2 \\
 -17
 \end{array}
 \end{array}$$

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

36

Is data being share?



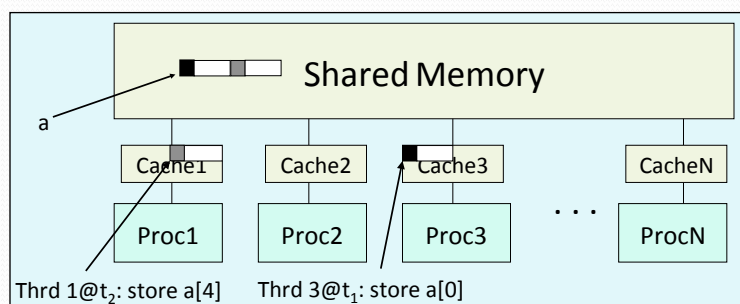
- Cache loads and stores work with 4-16 word long cache lines (64B for Intel)
 - If two threads are writing to the same cache line, conflicts occurs
- Even if the address differs, one will still suffer **performance penalty**
- This is called **false sharing**

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

37

Solution for False Sharing



- Be aware of the cache line sizes for a platform
- Avoid accessing the same cache line from different threads

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

38

Example 6: Eliminate False Sharing

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double pi; int nthreads;
    step = 1.0/(double) num_steps;
    #pragma omp parallel num_threads(NUM_THREADS)
    {
        int i, id = omp_get_thread_num(); double x, local_sum=0.0;
        int nthrds = omp_get_num_threads();
        for (i=id; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            local_sum += 4.0/(1.0+x*x);
        }
        #pragma omp critical
        pi += local_sum * step;
    }
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds



1.15 seconds

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

39

Example 6: Eliminate False Sharing

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double pi, x, local_sum; int nthreads;
    step = 1.0/(double) num_steps;
    #pragma omp parallel num_threads(NUM_THREADS) private(x, local_sum)
    {
        int i, id = omp_get_thread_num(); local_sum=0.0;
        int nthrds = omp_get_num_threads();
        for (i=id; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            local_sum += 4.0/(1.0+x*x);
        }
        #pragma omp critical
        pi += local_sum * step;
    }
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds



1.15 seconds

Great!
But is this the
best we can
do?

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

40

Example 7: Sequential Optimization

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double pi, x, local_sum; int nthreads;
    step = 1.0/(double) num_steps;
    #pragma omp parallel num_threads(NUM_THREADS) private(x, local_sum)
    {
        int i, id = omp_get_thread_num(); local_sum=0.0;
        int nthrds = omp_get_num_threads();
        double hoop = nthrds*step;
        for (i=id, x=(i+0.5)*step; i< num_steps; i=i+nthrds) {
            x += hoop;
            local_sum += 4.0/(1.0+x*x);
        }
        #pragma omp critical
        pi += local_sum * step;
    }
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds

0.46 seconds

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code? Carnegie Mellon University (c) 2012-2014 41

Outline

- Mutlicore Shared-Memory Platforms
- OpenMP – API for Shared-Memory Multi-processing
 - Overview
 - Example: numerical integration of Pi
 - Shared variable synchronization
 - **SPMD vs worksharing**
 - Data environment options
 - Advanced worksharing options
 - Synchronization options
- How is This Relevant to Writing Fast Code?

SPMD vs. Worksharing

- A parallel construct by itself creates an SPMD
 - “Single Program Multiple Data”
 - Programmer must explicitly specify what each thread must do differently
 - The division of work is hard-coded in the program
- Opportunity:
 - Many parallel regions are loops
- Question:
 - Can we make it easier to parallelize these loops?

OpenMP: Parallel For

- One of the worksharing constructs OpenMP provide is “**omp for**”

Sequential code

```
for(i=0;i<N;i++) { a[i] = a[i] + b[i];}
```

OpenMP “**parallel**”
region

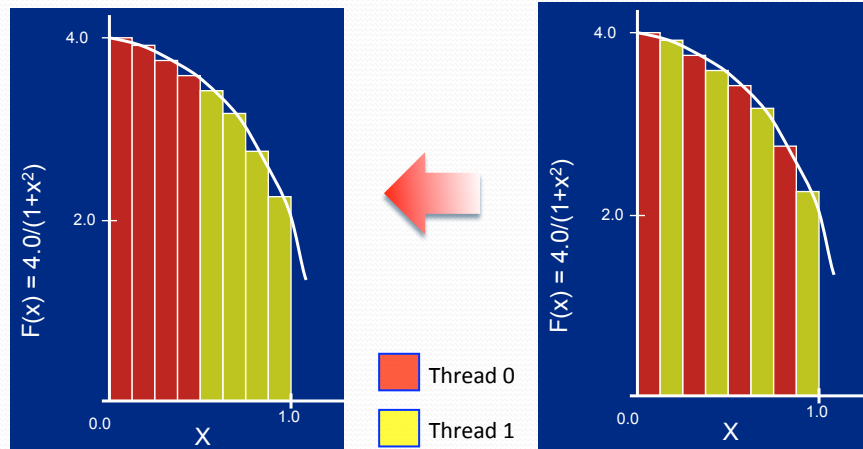
```
#pragma omp parallel
{
    int id, i, Nthrds, istart, iend;
    id = omp_get_thread_num();
    Nthrds = omp_get_num_threads();
    istart = id * N / Nthrds;
    iend = (id+1) * N / Nthrds;
    if (id == Nthrds-1) iend = N;
    for(i=istart;i<iend;i++) { a[i] = a[i] + b[i];}
}
```

OpenMP “**parallel**”
region and a
worksharing “**for**”
construct

```
#pragma omp parallel
#pragma omp for
for(i=0;i<N;i++) { a[i] = a[i] + b[i];}
```

Based on slides from Mattson et al, A “Hands-on” Introduction to OpenMP

Other Concurrency Opportunities



Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

45

Example 8: Parallel For and Reduction

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double pi, x, sum = 0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel private(x)
    {
        #pragma omp for reduction(+:sum)
        for (i=0; i< num_steps; i++){
            x = (i+0.5)*step;
            sum = sum + 4.0/(1.0+x*x);
        }
    }
    pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds

1.22 seconds

Reduction is more scalable than critical sections.

Loops must not have loop carry dependency.
Limitation of compiler technology.

"i" is private by default

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

46

Combined Parallel/Worksharing Construct

- OpenMP shortcut: Put the “parallel” and the worksharing directive on the same line

```
double res[MAX]; int i;
#pragma omp parallel
{
    #pragma omp for
    for (i=0; i< MAX; i++) {
        res[i] = huge();
    }
}
```



```
double res[MAX]; int i;
#pragma omp parallel for
for (i=0; i< MAX; i++) {
    res[i] = huge();
}
```

Based on slides from Mattson et al, A “Hands-on” Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

47

OpenMP: Working With Loops

- Basic approach
 - Find compute intensive loops
 - Make the loop iterations independent .. So they can safely execute in any order without loop-carried dependencies
 - Place the appropriate OpenMP directive and test

```
double hoop = nthrds*step;
for (i=id, x=(i+0.5)*step; i< num_steps; i=i+nthrds) {
    x += hoop;
    sum = sum + 4.0/(1.0+x*x);
}
```



```
for (i=0; i< num_steps; i++){
    x = (i+0.5)*step;
    sum = sum + 4.0/(1.0+x*x);
}
```

Based on slides from Mattson et al, A “Hands-on” Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

48

OpenMP: Reductions

```
for (i=0; i< num_steps; i++){
    x = (i+0.5)*step;
    sum = sum + 4.0/(1.0+x*x);
}
```



```
#pragma omp for reduction(+:sum)
for (i=0; i< num_steps; i++){
    x = (i+0.5)*step;
    sum = sum + 4.0/(1.0+x*x);
}
```

reduction (op : list)

1. A local copy of each list variable is made and initialized depending on the "op" (e.g. 0 for "+")
2. Updates occur on the local copy
3. Local copies are reduced into a single value and combined with the original global value

- Accumulating values into a single variable (sum) creates true dependence between loop iterations that can't be trivially removed
- This is a very common situation ... it is called a "reduction".
- Support for reduction operations is included in most parallel programming environments.

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

OpenMP: Reduction Operators

- Many different associative operands can be used with reduction:
- Initial values are the ones that make sense mathematically.

Operator	Initial value
+	0
*	1
-	0

Operator	Initial value
&	~0
	0
^	0
&&	1
	0

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

Example 8: Parallel For and Reduction

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;

int main ()
{
    int i; double pi, x, sum = 0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel for private(x) reduction(+:sum)
    for (i=0; i< num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds



1.22 seconds

Advantage: we created a parallel program without changing any code and by adding 2 simple lines of text!

Disadvantage:

- 1) Lot's of pitfalls if you don't understand system architecture
- 2) The basic result may not be the fastest one can do

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

Outline

- Multicore Shared-Memory Platforms
- OpenMP – API for Shared-Memory Multi-processing
 - Overview
 - Example: numerical integration of Pi
 - Shared variable synchronization
 - SPMD vs worksharing
 - **Data environment options**
 - Advanced worksharing options
 - Synchronization options
- How is This Relevant to Writing Fast Code?

Data Environment in OpenMP

- Shared Memory programming model:
 - Most variables are shared by default
- Global variables are **SHARED** among threads
 - File scope variables, static
 - Dynamically allocated memory (ALLOCATE, malloc, new)
- But not everything is shared...
 - Functions called from parallel regions are **PRIVATE**
 - Automatic variables within a statement block are **PRIVATE**.

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

53


Example: Data Sharing

```
double A[10];
int main() {
    int index[10];
    #pragma omp parallel
        work(index);
    printf("%d\n", index[0]);
}
```

```
extern double A[10];
void work(int *index) {
    double temp[10];
    static int count;
    ...
}
```

A, index and count are shared by all threads.

temp is local to each thread



Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

54

Changing Storage Attributes

- One can selectively change storage attributes for constructs using the following clauses:
 - SHARED
 - PRIVATE
 - FIRSTPRIVATE
- The final value of a private inside a parallel loop can be transmitted to the shared variable outside the loop with:
 - LASTPRIVATE

Example 9: firstprivate example

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double pi, x, local_sum = 0.0; int nthreads;
    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel private(x) firstprivate(local_sum)
    {
        int i, id = omp_get_thread_num(); local_sum=0.0;
        int nthrds = omp_get_num_threads();
        double hoop = nthrds*step;
        for (i=id, x=(i+0.5)*step; i< num_steps; i=i+nthrds) {
            x += hoop;
            local_sum += 4.0/(1.0+x*x);
        }
        #pragma omp critical
        pi += local_sum * step;
    }
    printf("pi = %f \n", pi);
    return 0;
}
```

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

Outline

- Multicore Shared-Memory Platforms
- OpenMP – API for Shared-Memory Multi-processing
 - Overview
 - Example: numerical integration of Pi
 - Shared variable synchronization
 - SPMD vs worksharing
 - Data environment options
 - **Advanced worksharing options**
 - Synchronization options
- How is This Relevant to Writing Fast Code?

The Scheduling Cause

- The schedule clause affects how loop iterations are mapped onto threads
 - `schedule(static [,chunk])`
 - Deal-out blocks of iterations of size “chunk” to each thread.
 - `schedule(dynamic[,chunk])`
 - Each thread grabs “chunk” iterations off a queue until all iterations have been handled.
 - `schedule(guided[,chunk])`
 - Threads dynamically grab blocks of iterations. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds.
 - `schedule(runtime)`
 - Schedule and chunk size taken from the OMP_SCHEDULE environment variable (or the runtime library ... for OpenMP 3.0).

Based on slides from Mattson et al, A “Hands-on” Introduction to OpenMP

The Scheduling Cause

Schedule Clause	When To Use	
STATIC	Pre-determined and predictable by the programmer	Least work at runtime : scheduling done at compile-time
DYNAMIC	Unpredictable, highly variable work per iteration	Most work at runtime : complex scheduling logic used at run-time
GUIDED	Special case of dynamic to reduce scheduling overhead	

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

59

Example 9: Scheduling

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;

int main ()
{
    int i; double pi, x, sum = 0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel for private(x) reduction(+:sum) schedule(static,
    10000000)
    for (i=0; i< num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds

1.31 seconds

Based on slides from Mattson et al, A "Hands-on" Introduction to OpenMP

18-645 – How to Write Fast Code?

Carnegie Mellon University (c) 2012-2014

60

Outline

- Mutlicore Shared-Memory Platforms
- OpenMP – API for Shared-Memory Multi-processing
 - Overview
 - Example: numerical integration of Pi
 - Shared variable synchronization
 - SPMD vs worksharing
 - Data environment options
 - Advanced worksharing options
 - **Synchronization options**
- How is This Relevant to Writing Fast Code?

Synchronization: ordered

- The **ordered** region executes in the sequential order
- Important for some scientific code and optimization code
 - Order of reduction may cause rounding differences

```
#pragma omp parallel private (tmp)
#pragma omp for ordered reduction(+:res)
    for (I=0; I<N; I++){
        tmp = NEAT_STUFF(I);
        #pragma ordered
        res += consum(tmp);
    }
```

Synchronization Barrier

- **Barrier:** Each thread waits until all threads arrive

```
#pragma omp parallel shared (A, B, C) private(id)
{
    id=omp_get_thread_num();
    A[id] = big_calc1(id);
    #pragma omp barrier
    #pragma omp for
    for(i=0;i<N;i++){C[i]=big_calc3(i,A);}
    #pragma omp for nowait
    for(i=0;i<N;i++){ B[i]=big_calc2(C, i); }
    A[id] = big_calc4(id);
}
```

implicit barrier at the end of a for worksharing construct

No implicit barrier due to **nowait**

Implicit barrier at the end of a parallel region

“Single” Worksharing Construct

- The **single** construct denotes a block of code that is executed by only one thread (not necessarily the master thread).
- A barrier is implied at the end of the single block
 - can remove the barrier with a **nowait** clause

```
#pragma omp parallel
{
    do_many_things();
    #pragma omp single
    { exchange_boundaries(); }
    do_many_other_things();
}
```

```
#pragma omp parallel
{
    do_many_things();
    #pragma omp single nowait
    { exchange_boundaries(); }
    do_many_other_things();
}
```


Nested Parallelism

- **Q:** Is nested parallel possible with OpenMP?
- **A:** Yes. But be sure to understand why you want to use it.

```
#include <omp.h>
#include <stdio.h>
void report_num_threads(int level) {
    #pragma omp single
    { printf("L %d: # threads in team %d\n", level, omp_get_num_threads()); }
}
int main() {
    omp_set_dynamic(0);
    #pragma omp parallel num_threads(2)
    {
        report_num_threads(1);
        #pragma omp parallel num_threads(2)
        {
            report_num_threads(2);
            #pragma omp parallel num_threads(2)
            {
                report_num_threads(3);
            }
        }
    }
    return(0);
}
```

Source from: <http://download.oracle.com/docs/cd/E19205-01/819-5270/aewbc/index.html>

Nested Parallelism

- Compiling and running this program with nested parallelism enabled produces the following (sorted) output:

```
$ export OMP_NESTED=TRUE
$ ./experimentN
L 1: # threads in team 2
L 2: # threads in team 2
L 2: # threads in team 2
L 3: # threads in team 2
L 3: # threads in team 2
L 3: # threads in team 2
L 3: # threads in team 2
```

```
$ export OMP_NESTED=FALSE
$ ./experimentN
L 1: # threads in team 2
L 2: # threads in team 1
L 3: # threads in team 1
L 2: # threads in team 1
L 3: # threads in team 1
```


OpenMP Environment Variables

- OMP_SCHEDULE=algorithm
 - dynamic[, n]
 - guided[, n]
 - Runtime
 - **static[, n]**
- OMP_NUM_THREADS=num
- OMP_NESTED=TRUE|FALSE
- OMP_DYNAMIC=TRUE|FALSE

Outline

- Mutlicore Shared-Memory Platforms
- OpenMP – API for Shared-Memory Multi-processing
 - Overview
 - Example: numerical integration of Pi
 - Shared variable synchronization
 - SPMD vs worksharing
 - Data environment options
 - Advanced worksharing options
 - Synchronization options
- **How is This Relevant to Writing Fast Code?**

k-means Algorithm Concurrency

1. **Initialization:** Randomly select k cluster centers
 2. **Expectation:** Assign closest center to each data point
 - N (independent)
 - k (min reduction)
 - D (sum reduction)
 3. **Maximization:** Update centers based on assignments
 - D (independent)
 - N (Histogram computation into k bins)
 4. **Evaluate:** Re-iterate steps 2-3 until convergence
- 

Project 1: Search Among Alternatives

- What are the implementation **alternatives**?
 - Different mapping of **application concurrency** to **platform parallelism**
- The **search** process
 - More complex than one **application concurrency** to one **platform parallelism**
 - May want to sequentialize some operations:
 - Some parallel operations are as “**work-efficient**” as sequential operations
 - Reduction – sequential: $O(N)$, Parallel: $O(N \log N)$
 - One level of concurrency could map to multiple levels of parallelism

How is this relevant to writing fast code?

Fast Platforms

- Multicore platforms
- Manycore platforms
- Cloud platforms



Good Techniques

- Data structures
- Algorithms
- Software Architecture

- This lecture: **An abstraction for multicore parallel programming**
- **Abstractions:**
 - Help establish a mental model for the programmers
 - Make it easier to write parallel code
 - **Performance depend on deep understanding of the implementation platform**

Can You Answer These Questions Now?

- What are the exploitable levels of parallelism in a multicore processor?
- What is SPMD? And how to use OpenMP to do SPMD?
- What's the difference between "critical" and "atomic"?
- How to reduce synchronization cost and avoid "false sharing"?
- What are the scheduling, reduction, data sharing, and synchronization options for OpenMP?
- How is this relevant to writing fast code?