

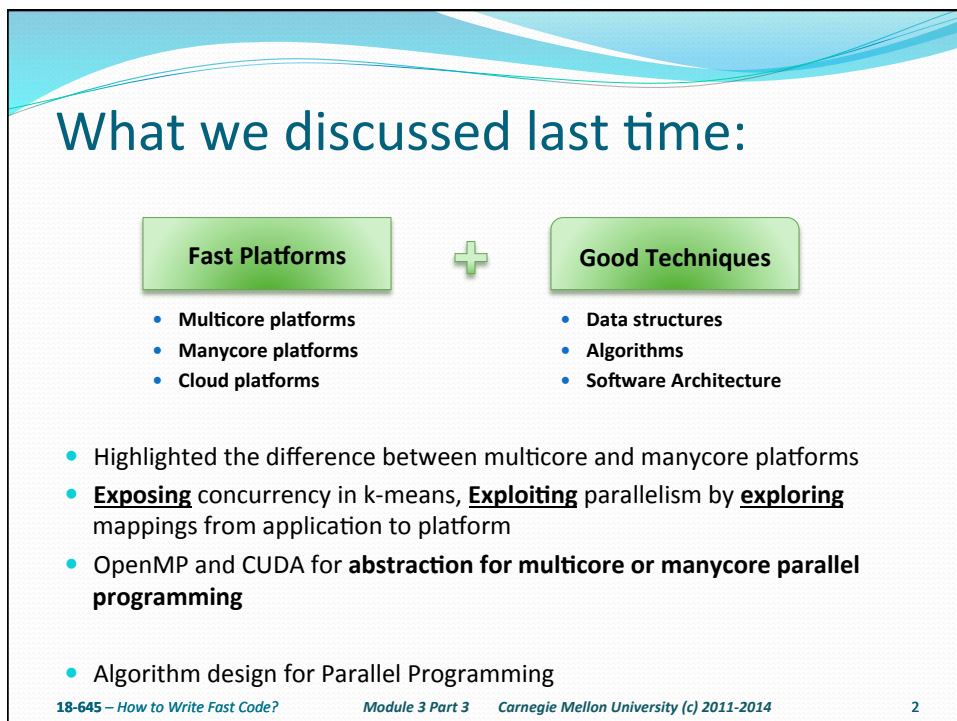
Module 3 Part 3

Application Design for Manycore Data Parallel Algorithms

Carnegie Mellon University
18-645

Jike Chong
Ian Lane

18-645 – How to Write Fast Code? 1



What we discussed last time:

Fast Platforms

- Multicore platforms
- Manycore platforms
- Cloud platforms

+

Good Techniques

- Data structures
- Algorithms
- Software Architecture

- Highlighted the difference between multicore and manycore platforms
- **Exposing** concurrency in k-means, **Exploiting** parallelism by **exploring** mappings from application to platform
- OpenMP and CUDA for **abstraction for multicore or manycore parallel programming**
- Algorithm design for Parallel Programming

18-645 – How to Write Fast Code? Module 3 Part 3 Carnegie Mellon University (c) 2011-2014 2

Answers you should know after this...

- What are the important properties of a Map function?
- What are the important properties of a Reduce function?
- What are the important properties of a Scan function?


- How to compact an array in a data-parallel way?
- How to find unique elements in an array in a data-parallel way?

Outline

- Algorithm Optimization
 - Data Parallel Algorithms
 - Map, Reduce, and Scan
 - Common Compositions of Data Parallel Algorithms
 - Compact, Find Unique, Building a Flag Array

Data Parallel Algorithms - Map

- **Map :**
A function that applies a given function to each element of a list, and returning a list of results



- Two important properties:
- **Side-effect free:**
Only returning a value, no modifications of state with the rest of the application
- **Independent:**
Has an independent piece of work, where its input does not depend on another function


Data Parallel Algorithms - Reduce

- **Reduce:**
A function that takes in a list of objects and builds up a return value
- Important properties for parallel reduction:
 - Associativity: $a+(b+c) == (a+b)+c$
 - Allows elements to be reduced in parallel in a “tree”
- In CUDA, the synchronization has to be managed by the programmer

$$\begin{aligned}
 &a+b+c+d+e+f+g+h \\
 &= ((a+b)+(c+d)) + ((e+f)+(g+h)) \\
 &= (a+b+c+d) + (e+f+g+h)
 \end{aligned}$$

How Best to Implement Reduce?

- What is an issue with this approach?




```
unsigned int t = threadIdx.x;
for (unsigned int stride = 1;
     stride < blockDim.x; stride *= 2)
{
    __syncthreads();
    if (t % (2*stride) == 0)
        partialSum[t] += partialSum[t+stride];
}
```

- No more than half of threads will be executing at any time

How Best to Implement Reduce?

- What about this approach?




```
unsigned int t = threadIdx.x;
for (unsigned int stride = blockDim.x;
     stride > 1; stride >> 1)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] += partialSum[t+stride];
}
```

- Minimize branch divergence

Elimination of __syncthreads()

```
__device__ void sum(float* g_data, float* g_o_data){
    unsigned int tid = threadIdx.x;
    extern __shared__ float s_data[];
    // Assign initial value
    s_data[tid] = g_data[...];
    __syncthreads();
    if (tid < 128)
        s_data[tid] += s_data[tid + 128];
    __syncthreads();
    if (tid < 64)
        s_data[tid] += s_data[tid + 64];
    __syncthreads();
    if (tid < 32) {
        // No __syncthreads() 32 threads in each
        // warp execute in lock-step with each other
        volatile float* s_ptr = s_data;
        s_ptr[tid] += s_ptr[tid + 32];
        s_ptr[tid] += s_ptr[tid + 16];
        s_ptr[tid] += s_ptr[tid + 8];
        s_ptr[tid] += s_ptr[tid + 4];
        s_ptr[tid] += s_ptr[tid + 2];
        s_ptr[tid] += s_ptr[tid + 1];
    }
    // Write result for this thread block to global memory
    if (tid == 0)
        g_o_data[blockIdx.x] = s_data[0];
}
```



18-645 – How to Write Fast Code?

Module 3 Part 3

Carnegie Mellon University (c) 2011-2014

9

Data Parallel Algorithms - Scan

- **Scan (prefix-sum):**
Takes a binary associative operator \oplus with identity I , and an array of n elements $[a_0, a_1, \dots, a_{n-1}]$ and returns the ordered set $[I, a_0, (a_0 \oplus a_1), \dots, (a_0 \oplus a_1 \oplus \dots \oplus a_{n-2})]$.
- **Example:**
if \oplus is addition, then scan on the set $[3, 1, 7, 0, 4, 1, 6, 3]$
returns the set $[0, 3, 4, 11, 11, 15, 16, 22]$
- **How fast can we do that?**

18-645 – How to Write Fast Code?

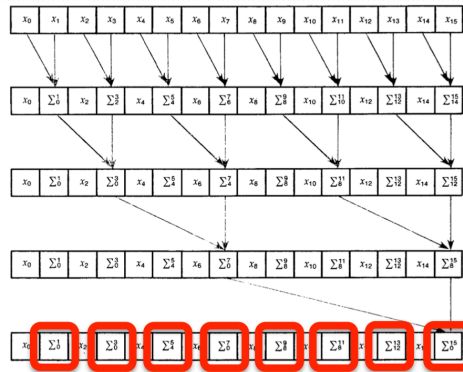
Module 3 Part 3

Carnegie Mellon University (c) 2011-2014

10

To Implement Scan – Revisit Reduce

- Any techniques for creating for an $O(\log N)$?



Sequential Reduction


```
sum[1] = v[1]
For k = 2 to N
    sum[k] = sum[k-1] + v[k]
```

Parallel Reduction

How do you do a scan?

Scan Techniques

- Note that in the vector reduction, at least half of the processors were idle in any step. Let's have them compute something!



Parallel Reduction

```
for j = 1 to log2(n),
    for all k in parallel,
        if ((k+1) mod 2^j) == 0,
            x[k] = x[k - 2^{j-1}] + x[k]
```

Parallel Prefix Sum

```
for j = 1 to log2(n),
    for all k in parallel,
        if k >= 2^j,
            x[k] = x[k - 2^{j-1}] + x[k]
```

Scan Libraries

- Like **sort**, there exist many optimizations for **scan**
- Expert in the area: Professor Guy Blelloch
 - 15-499: Parallel Algorithms
- For writing fast scan implementations – use the **Thrust library**
 - Thrust library: C++ template library for CUDA
 - Now part of CUDA 4.0

```
#include <thrust/scan.h>

int data[6] = {1, 0, 2, 2, 1, 3};
thrust::exclusive_scan(data, data + 6, data); // in-place scan

// data is now {0, 1, 1, 3, 5, 6}
```

<http://code.google.com/p/thrust/wiki/QuickStartGuide#Prefix-Sums>


Outline

- Algorithm Optimization
 - Data Parallel Algorithms
 - Map, Reduce, and Scan
 - Common Compositions of Data Parallel Algorithms
 - Compact, Find Unique, Building a Flag Array

Data Parallel Algorithms - Compact

- Compaction:**

Removing elements from an array - take in an array, and produce an shorter array



- How do we perform removal in parallel?

Data Parallel Algorithms - Compact

- Compaction:**

Removing elements from an array - take in an array, and produce an shorter array

- How do we perform removal in parallel?

Removing: [1] [5] [6] [8]

A	B	C	D	E	F	G	H	I	J
---	---	---	---	---	---	---	---	---	---

- Map** – create flags (“1” keep, “0” remove)

Flags

1	0	1	1	1	0	0	1	0	1
---	---	---	---	---	---	---	---	---	---


- Scan** – compute index

scanIdx

0	1	1	2	3	3	3	4	4	5
---	---	---	---	---	---	---	---	---	---

- Map** – copy to new array


```
if (flag[i] == 1){
    dst[ scanIdx[i] ] = src[i];
}
```



Data Parallel Algorithms - FindUniq

- FindUniq:**


Removing duplicates from an array – take in an array, produces a equal or smaller set of unique values



- How do we perform “find unique” in parallel?

- How do we “find unique” sequentially?

- Sort
- Iterate through and copy



18-645 – How to Write Fast Code?

Module 3 Part 3


Carnegie Mellon University (c) 2011-2014

17

Data Parallel Algorithms - FindUniq

- FindUniq:**

Removing duplicates from an array – take in an array, produces a equal or shorter array




- How do we perform “find unique” in parallel?


- Sort**




- Map** – flag when i^{th} and $(i-1)^{\text{th}}$ element differ
[0] = 1



- Scan** – create compaction index



- Map** – copy to new array



18-645 – How to Write Fast Code?


Module 3 Part 3

Carnegie Mellon University (c) 2011-2014

18

Find Unique – Special case

Traditional Approach



0.349 seconds

- Special case:
What if we know all the possible values the elements can take – such as the 26 letters of the alphabet.

- Sorting is the most expensive step
- How can we avoid sorting?


18-645 – How to Write Fast Code?

Module 3 Part 3 Carnegie Mellon University (c) 2011-2014

19

Find Unique – Special case

Traditional Approach



0.349 seconds

- Special case:
What if we know all the possible values the elements can take – such as the 26 letters of the alphabet.


- Sorting is the most expensive step
- How can we avoid sorting?

Setup a hash table, or lookup table

A	B	C	...	I	J	...	M	...	P	...	S
---	---	---	-----	---	---	-----	---	-----	---	-----	---

0	0	0	...	1	0	...	1	...	1	...	1
---	---	---	-----	---	---	-----	---	-----	---	-----	---

Alternative Approach



0.055 seconds

Jike Chong, Ekaterina Gonina, Kurt Keutzer, "Efficient Automatic Speech Recognition on the GPU", Chapter in GPU Computing Gems Emerald Edition, Morgan Kaufmann, Vol. 1, February 9, 2011.

18-645 – How to Write Fast Code?

Module 3 Part 3 Carnegie Mellon University (c) 2011-2014


20

Populating a Hash Flag Array

- Hash insertion:
 - Leverage the semantics of conflicting writes for non-atomic memory accesses
 - At least one conflicting write to a device memory location is guaranteed to succeed
 - Order of execution is undefined

NVIDIA CUDA C Programming Guide Version 4.0, Chapter 4.2 (page 86)

- For setting flags, map the hash insertion to threads
 - success of insertion in any order can achieve the goal



Alternative Approach

Hash insertion

Hash write (0.030)

Duplicate Removal

Unique-index
Prefix-scan (0.020)

Unique-list
Gathering (0.005)

0.055 seconds

How is this relevant to writing fast code?

Fast Platforms

- Multicore platforms
- Manycore platforms
- Cloud platforms



Good Techniques

- Data structures
- Algorithms
- Software Architecture

- Introduced the manycore platform HW and SW mental models
- Introduced the terminologies for you to start FLIRTING with the technology
- Introduced design trade-offs in data structures
- Introduced parallel algorithms
- Next lectures: Focus on software architecture